

Designing Human–AI Co-Creation Systems for Movie Commentary Videos

BAIXUE WU

1 Introduction

Movie commentary videos have become a widely consumed digital media format in which creators summarize, interpret, or critique films through narrative explanation paired with selected audiovisual clips [26]. These videos span professional analyses that examine cinematic techniques and directorial style, as well as popular online formats where creators humorously recount movie plots for broad audiences—for example, the film-review channels widely found on YouTube and Bilibili [9, 14]. They serve as an accessible bridge between cinematic works and viewers, facilitating the reinterpretation of complex plots, aesthetics, and emotions through engaging narrative framing. As social media platforms continue to grow, such videos have become a prominent medium for film communication—enabling enthusiasts, educators, and critics to share insights and promote participatory discussion [29, 46]. The demand for producing high-quality movie commentary videos is substantial, yet the workflow remains labor-intensive, requiring creators to script, select, and align audiovisual materials while maintaining narrative coherence and emotional rhythm.

However, producing a movie commentary video remains a highly demanding and time-consuming task. Creators must first identify visual materials that correspond to the narrative script and ensure that the selected clips align with the intended emotional tone and pacing of the video [7]. When assembling these clips, creators must maintain narrative coherence and avoid abrupt transitions or excessive jump cuts that could disrupt the storytelling flow [19]. Beyond visual editing, creators often need to supplement contextual knowledge referenced in the script—such as historical background or production details—that is not directly apparent in the footage [20]. Finally, additional tasks such as voiceover recording, background music design, and subtitle synchronization demand further technical effort [2], rendering the overall production process cognitively intensive and creatively taxing.

Existing video editing software offers various functions that help creators reduce the technical burden of editing. Professional tools such as *Adobe Premiere Pro* provide features like automatic audio–video synchronization, subtitle generation [1], and extensive template libraries, while lightweight platforms like *CapCut* support quick transitions and visual effects through accessible presets [6]. However, these tools still require users to navigate complex interfaces and follow predefined editing conventions to achieve coherent results. The emergence of AI-driven editing technologies has further advanced automation by supporting tasks such as footage retrieval, knowledge integration, and voice synthesis [10, 13, 36]. Yet these intelligent capabilities remain fragmented across separate tools and APIs, requiring creators to manually integrate multiple systems within their workflow. A unified, human-centered framework that integrates these AI capabilities into a coherent production pipeline remains largely absent.

To address this gap, the proposed research aims to develop **Movie-Explain-RAG**, a human–AI co-creation system intended to transform commentary scripts into narratively coherent and emotionally expressive video drafts. Rather than pursuing full automation, the system will collaborate with creators—providing suggestions, surfacing alternatives, and adapting to human decisions within a unified workflow that will integrate semantic understanding, knowledge augmentation, and temporal organization.

Conceptually, **Movie-Explain-RAG** will operate through four coordinated stages: (1) it will interpret the input script by segmenting it into narrative and affective units; (2) it will identify

missing or weakly aligned visual evidence and retrieve candidate clips or contextual materials using retrieval-augmented generation (RAG) [24]; (3) it will assemble selected materials into a structurally coherent rough cut guided by temporal and cinematic pacing heuristics; and (4) it will support iterative human–AI refinement by enriching the timeline with subtitles, narration, audio design, and other presentation elements that creators can fine-tune in real time. By combining multimodal retrieval, knowledge-enhanced generation, and mixed-initiative editing, the project will investigate how AI can function as an *interpretive collaborator*—helping creators produce commentary videos that are semantically grounded, emotionally resonant, and collaboratively authored.

To evaluate its feasibility and impact, a mixed-method research design will be employed, integrating quantitative and qualitative analyses. The first study (E1) will quantitatively assess alignment accuracy—semantic, affective, and temporal—using both automatic and expert-based metrics against retrieval-only and RATV baselines. The second study (E2) will qualitatively examine how creators will collaborate with the system throughout rough-cut assembly and A/V refinement, observing how AI-assisted retrieval, structural sequencing, subtitle generation, and audio alignment will influence workflow, decision-making, and perceived authorship through interaction logs and semi-structured interviews. Together, these studies will combine quantitative validation of multimodal coherence with qualitative insights into mixed-initiative creative practice, demonstrating how Movie- Explain-RAG will enhance both technical precision and human interpretability in script-to-video co-creation.

2 Related Work

2.1 User-generated Content

User-generated content (UGC) illustrates how audiences have shifted from passive spectators to active co-creators of media [4]. Early studies framed this shift as participatory culture or user-generated creativity, where individuals reinterpret existing media texts through expressive acts [8].

In HCI, a growing body of research has examined diverse manifestations of such re-creative behavior. Studies on fan-made videos and machinima show how audiences appropriate cinematic or game materials to convey commentary, humor, or critique—often constructing alternative storylines that challenge mainstream narratives [12]. Investigations into remix culture have further highlighted the social and technical infrastructures that enable such creativity [32], ranging from meme-editing platforms to large-scale collaborative editing communities that coordinate multi-user video assemblage [28, 42]. Amid this landscape, movie commentary videos stand out as a form of narrative reinterpretation—combining voiceovers, curated visuals, and subtitles [26]. Collectively, this body of work reveals a paradigm shift in media interaction—from the consumption of finished works to their reinterpretation and rewriting through digital tools.

2.2 Tools for Text-to-Video Editing

Text-to-video editing tools leverage language understanding to retrieve and assemble video clips from textual scripts. Early systems such as QuickCut [37] aligned narrated transcripts with footage to streamline rough-cut generation. Write-A-Video [38] extended this approach by segmenting scripts into semantic units and retrieving temporally coherent clips using visual–semantic embeddings.

Subsequent work further refined this paradigm. Transcript-to-Video [43] introduced a weakly supervised method for aligning textual cues with relevant clips while preserving temporal order. Shot Retrieval [44] employed sequence-level pretraining to better capture narrative flow and long-range dependencies. Lotus [5] examined how combining abstractive and extractive summarization can preserve emotional pacing and aesthetic continuity during video montage. Collectively, these systems reflect a shift from transcript-aligned editing toward semantic and stylistic assembly.

2.3 Retrieval-augmented Generation

Retrieval-Augmented Generation (RAG) combines external retrieval with generative modeling, enabling systems to ground outputs in relevant knowledge or context [24]. Originally developed for factual tasks such as QA and reasoning [15], RAG has since expanded into creative domains including scientific writing, code generation, and narrative design [3, 34, 41].

In the context of creative and multimodal content generation, RAG operates not only as a factual retrieval mechanism but also as a source of narrative and stylistic inspiration. It has been applied to story generation through text–image retrieval [30], visual design via exemplar-based suggestions [35], and music composition using audio references [23]. Beyond text and design, retrieval-augmented methods have supported video production in domains such as news and documentary editing [39]. Collectively, these applications demonstrate how RAG bridges knowledge grounding and creative expression, transforming retrieval into an interactive and interpretive process.

2.4 Human-AI Co-Creation

Human–AI co-creation describes interactive processes in which humans and intelligent systems jointly contribute to creative outcomes. Unlike fully automated systems, co-creative tools emphasize shared agency, wherein both human and AI shape the creative trajectory [48]. The rise of large language models has enabled new forms of collaboration across domains such as writing, education, and design, allowing systems to provide contextual suggestions, interpret intent, and adapt to user feedback [47].

In co-creative media authoring, human–AI collaboration has been examined across multiple modalities, including video, audio, and editing. In video creation, studies focus on video summarization and re-narrativization, where AI systems generate rough cuts or condensed storylines that users reorganize to refine emotional rhythm and narrative perspective [40]. In audio domains, collaboration centers on sound design, music composition, and voice or timbre generation. Here, AI assists in producing initial soundscapes, melodies, or vocal performances, while users adjust tone, tempo, and expressive nuance to match intended affect [16, 49]. In editing, research emphasizes multimodal collaborative platforms and cross-modal synchronization, integrating voiceover, subtitles, and visual rhythm to maintain temporal and aesthetic coherence [21]. Across these modalities, human–AI co-creation reframes media production as an iterative, multimodal dialogue between automation and human interpretation.

3 Method

3.1 Design Guidelines

This project adopts a **design-oriented, concept-driven HCI methodology** that employs system prototyping as a means to reason about human–AI creative collaboration [31, 50]. Rather than emphasizing algorithmic optimization, this approach focuses on generating design knowledge about how AI systems can support interpretive, affective, and co-creative practices in multimodal storytelling. Through this lens, the project develops design guidelines and an envisioned system concept to articulate how human and AI contributions may be orchestrated throughout the commentary production workflow.

To identify recurring practices, we reviewed 30 deep-analysis commentary channels and observed how creators highlight motifs, reveal foreshadowing, and explain directorial intent. We then examined challenges noted in multimodal alignment and narrative editing. Difficulties in synchronizing textual interpretation with visual affect and detecting narrative-relevant events informed DG1 [33, 45]. Evidence that viewer comprehension depends strongly on shot order and pacing




Fig. 1. Overview of the Movie-Explain-RAG system pipeline.

shaped DG2 [25]. Frequent reliance on cultural or historical context motivated DG3 [24]. Finally, research on co-creative editing emphasizing editable drafts and reversible refinement—rather than one-shot automation—informed DG4 [40].

We present the final set of design guidelines below. The two overarching goals of these guidelines are (1) maintaining semantic and emotional coherence between commentary and visual narrative structure, and (2) foregrounding creator agency through iterative co-creation.

DG1: Semantic and Affective Alignment. Grounded in research identifying the affective gap between textual interpretation and visual meaning [17, 45] and in techniques for detecting narrative-relevant anchor events in video [33], the system should align commentary not only with the semantic content of selected clips but also with their emotional tone and narrative emphasis.

DG2: Narrative Temporal Structure. Prior work shows that shot ordering and temporal pacing strongly influence narrative cognition and audience engagement [22, 25, 51]. Thus, the system should support structured sequencing (e.g., setup–revisit–reveal) to preserve coherence when explaining symbolic motifs or thematic arcs across scenes.

DG3: Knowledge-Augmented Commentary. Deep-analysis commentary frequently depends on contextual information such as cultural references, directorial intention, or production history. Drawing on retrieval-augmented generation for knowledge-intensive reasoning [24] and on agentic video interpretation systems [11], the system should surface relevant external sources to substantiate interpretive claims while preserving the creator’s authorial voice.

DG4: Human–AI Iterative Co-Creation. Reflecting co-creative interaction frameworks and video editing systems that support alternative drafts and reversible refinement [18, 21, 27, 40], the system should enable creators to iteratively adjust narrative emphasis, pacing, and tone rather than rely on one-shot automation.

3.2 System Pipeline

Building on the design-oriented principles described above, we will introduce **Movie-Explain-RAG** as a system concept intended to support co-creative movie commentary video production. Rather than automating the process end-to-end, the system will collaborate with the creator—proposing options, responding to decisions, and adapting to human input throughout a structured workflow. This workflow will transform the commentary script into a narratively coherent, affectively expressive, and interpretively enriched video output. Each stage will operationalize the four design goals—semantic alignment, temporal coherence, knowledge augmentation, and human–AI co-creation—within an integrated pipeline coordinated by a central controller that will maintain project memory and version history.

Stage 1 – Script Interpretation & Concept Structuring. The process will begin when the creator uploads or drafts a commentary script. The system will perform semantic and affective parsing, segmenting the text into narrative units and encoding them with multimodal embeddings to capture literal meaning and emotional tone (**DG1**). Each segment will be automatically annotated with visualizable concepts (characters, settings, symbolic motifs), forming anchor points for subsequent retrieval.

Creators will be able to review, correct, and extend these tags—clarifying misinterpretations, specifying emotional intentions, or highlighting interpretive emphasis. These human refinements will be incorporated into the shared representational space that will be used by subsequent modules.

Stage 2 – Content Gap Analysis & Retrieval Augmentation. The system will perform *coverage gap analysis* to detect where visual evidence is missing or weakly aligned with the script. For emotion-oriented gaps, the *Content Retrieval Module (CRM)* [43] will surface candidate scenery or “empty” shots to preserve affective continuity (**DG1**). For knowledge-oriented gaps, the *Retrieval-Augmented Generation (RAG)* [24] layer will retrieve contextual or cultural images and will generate short explanatory text when needed (**DG3**). Retrieved materials will be evaluated through a multi-dimensional coherence model.

During this stage, creators will primarily refine RAG-generated content—correcting misinterpretations, adjusting contextual images or brief explanations, and replacing suggestions that do not fit their intended interpretation. These human refinements will guide retrieval direction and will inform how RAG updates its priorities, supporting mixed-initiative meaning-making while leaving clip selection for the next stage.

Stage 3 – Temporal Composition & Structural Rough Cut. Using the approved clips, a *Temporal Consistency Module (TCM)* [43] will assemble an initial rough cut that emphasizes structural temporal coherence (**DG2**). Guided by cinematographic pacing heuristics, the module will estimate clip ordering, scene boundaries, and segment duration to preserve narrative continuity and emotional rhythm.

At this stage, the system will operate primarily on the visual timeline, generating a draft sequence without final subtitles or detailed audio mixing. Creators will review the rough cut, adjust pacing, refine transitions, replace mismatched clips, or restructure segments to better express their intended narrative emphasis. These edits will update timing estimates and will provide a stable foundation for subsequent A/V refinement.

Stage 4 – A/V Enrichment & Iterative Human-AI Refinement. Building on the rough cut from Stage 3, the system will enrich the timeline with additional audio-visual layers—generating or aligning subtitles, synchronizing voice-over narration, and proposing background music or sound design options that will match the commentary’s emotional and thematic contours. It will also be able to suggest overlays or emphasis cues for key interpretive points.

Creators will then fine-tune these elements by editing subtitles, adjusting micro-timing between speech and visuals, selecting or replacing music, and adding contextual inserts or emphasis effects. All edits will be stored as interaction traces, enabling the system to adapt its generative behavior and to offer more targeted suggestions over time. This mixed-initiative loop will operationalize **DG4**, ensuring that creative authority remains with the human while reducing manual effort.

3.3 Scenario-based Workflow

To illustrate how **Movie-Explain-RAG** will support the creative workflow of commentary video production, we will outline a scenario that depicts a typical interaction cycle between a future user and the envisioned system. This scenario will not describe a finalized interface, but rather the anticipated human-AI dynamics that the system will aim to enable.

Scenario. A film student, Mei, will plan to create a commentary video analyzing themes of nostalgia and emotional restraint in Wong Kar-wai’s *In the Mood for Love*. She will begin by uploading a draft script containing her preliminary interpretations.

First, the system will interpret the script by segmenting it into narrative and affective units and highlighting emotionally salient expressions such as “muted desire” and “rhythmic stillness.” It will generate initial concept tags—characters, settings, symbolic motifs—that will guide later retrieval. Mei will review these suggestions, will correct several misaligned tags, and will adjust the emphasis to better match her interpretive intentions.

Next, the system will analyze the script for visual or contextual gaps. It will surface candidate film clips through the *Content Retrieval Module (CRM)* and will provide contextual images or textual references via *Retrieval-Augmented Generation (RAG)*. These materials will be presented as suggestions that Mei will be able to accept, reject, or refine. She will examine the proposed clips, will replace one with a moment she finds more emotionally resonant, and will incorporate a retrieved archival still that will reinforce her argument.

Afterward, the system will begin constructing a structural rough cut. Using cinematographic pacing heuristics, the *Temporal Consistency Module (TCM)* will estimate clip ordering, scene boundaries, and approximate durations to maintain narrative continuity. Mei will preview this draft, will identify transitions that feel abrupt, will lengthen several dissolves, and will adjust the arrangement of segments to align the visual rhythm with her interpretive stance. The system will update the timeline based on these revisions.

Finally, the structurally coherent rough cut will be enriched with detailed audio–visual elements. The system will generate subtitle drafts, will synchronize Mei’s narration, and will suggest background music options that will reflect the film’s emotional contour. Mei will edit the subtitles for clarity, will modify several lines of narration, will experiment with the proposed scores, and will fine-tune micro-timing across the final timeline. Throughout this process, the system will record her edits as *interaction traces*, enabling future suggestions to better reflect her stylistic and interpretive preferences.

4 Evaluation

To assess its feasibility and potential impact, the evaluation will be structured around a two-stage mixed-method strategy that will combine computational metrics with human-centered analysis. The first study (**E1**) will quantitatively examine the system’s ability to achieve semantic, affective, and temporal alignment. The second study (**E2**) will qualitatively explore how users will experience and interpret human–AI co-creation—particularly during structural rough-cut editing and A/V refinement—providing insights into usability, creative control, and perceived authorship.

4.1 E1: Technical Validity

Goal and Hypotheses. The first study will seek to quantitatively examine how effectively the proposed system will align commentary text with relevant video materials. Three hypotheses will be tested: **H1 (Semantic)**—the proposed method will achieve stronger semantic alignment than retrieval-only baselines; **H2 (Affective)**—it will better match affective tone and visual mood; and **H3 (Temporal Fit)**—it will achieve lower boundary misalignment between narration and clip transitions.

Data and Conditions. We will sample 30–50 commentary units from the envisioned scenario and several representative films. Each unit will contain five candidate clips that will be retrieved from an open media archive. Four conditions will be compared: (1) *Ours*—the full Movie-Explain-RAG

pipeline (CRM + RAG + TCM); (2) *Baseline-A*—a retrieval-only condition adapted from prior text-to-video retrieval frameworks such as *Write-A-Video* and *QuickCut*, which will rely solely on semantic similarity for clip selection without temporal or knowledge modeling [37, 38]; (3) *Baseline-B*—a stronger retrieval baseline modeled after the *Transcript-to-Video* system [43], which will implement the RATV framework combining a Content Retrieval Module (CRM) and a Temporal Consistency Module (TCM); and (4) an optional *Human-Upper* reference that will be curated by an experienced editor.

Metrics and Analysis. **Automatic metrics:** (i) *SemAlign*—CLIP-based text–video similarity and Recall@K, following prior work on text–video retrieval and clip sequencing [33, 43]; (ii) *AfAlign*—text–clip affective agreement (F1 or cosine similarity) based on affective video modeling and emotion-recognition frameworks [17, 45]; (iii) *Coverage*—keyword/entity overlap ratio for contextual retrieval consistency [24]; (iv) *TemporalFit*—timing alignment and optical-flow smoothness for temporal coherence, which will extend the temporal-consistency modeling in RATV [43, 51]. **Human evaluation:** Approximately fifteen expert raters will assess the *semantic*, *affective*, and *temporal* alignment quality of system outputs on 1–7 Likert scales, following established co-creative evaluation practices [18, 21, 40]. The purpose will be to validate the correspondence between automatic metrics and expert perception rather than to evaluate user experience. Inter-rater reliability will be measured using Krippendorff’s α , and correlation with automatic metrics (Spearman’s ρ) will establish convergent validity.

Expected Outcomes. It is expected that Movie-Explain-RAG will outperform retrieval-only baselines on semantic and affective alignment, while achieving comparable temporal accuracy to the human reference. These anticipated results will provide preliminary quantitative evidence of alignment accuracy, forming the empirical foundation for the subsequent user-centered study.

4.2 E2: User Study

Goal and Research Questions. Whereas **E1** will address quantitative correctness, the second study will focus on exploring how creators will perceive, interpret, and collaborate with the system in practice—particularly during structural rough-cut editing and A/V enrichment. Rather than merely evaluating usability, this study will aim to uncover the mechanisms of human–AI collaboration in creative work. It will investigate how creators will construct trust, interpretive agency, and affective intention when working with AI-assisted retrieval, rough-cut assembly, and audio–visual refinement.

Research Questions. This study will address: (1) how AI-assisted retrieval, rough-cut assembly, and A/V suggestions will influence users’ creative decisions and sense of agency; (2) how creators will negotiate the balance between automation and personal expression; and (3) how human–AI co-creation will shape trust, interpretive meaning, and perceived authorship.

Participants and Procedure. We will plan to recruit 12–16 participants (students or amateur creators with basic video-editing experience). Each participant will complete two comparable tasks: (1) *Ours*—using Movie-Explain-RAG to co-create a commentary video; and (2) *Baseline (NLE)*—using a standard non-linear editing tool with manual retrieval and timeline assembly, following conventional workflows that will be described in prior interactive video editing systems [37]. Each session will last approximately 45–60 minutes. All interactions will be screen-recorded and logged (retrieval actions, clip selections, timeline edits, subtitle adjustments, A/V refinements). After each session, participants will join a 10–15 minute semi-structured interview that will allow them to reflect on usability, creative control, interpretive influence, and cognitive effort.

Data Collection and Analysis. The study will collect and analyze: (1) system interaction logs (e.g., accepted AI suggestions, exploration patterns, fine-grained A/V adjustments); (2) reflective self-reports on perceived creativity, trust, and control; and (3) interview transcripts, which will be analyzed via thematic coding. Stimulated recall sessions—where participants will revisit key editing moments—will be used to uncover how AI suggestions will influence narrative, emotional, or stylistic choices. Identified themes will be organized under categories such as *creative agency*, *transparency and trust*, and *interpretive collaboration*.

Expected Outcomes. It is anticipated that participants will describe Movie-Explain-RAG as a supportive “co-author” that will enhance ideation and interpretive richness without diminishing creative ownership. The qualitative findings will be expected to complement the quantitative evidence from **E1**, collectively illustrating how human–AI collaboration will function as a *mixed-initiative creative process* in commentary video production. Whereas **E1** will ground the system’s multimodal coherence, **E2** will contextualize these metrics within lived creative practice—together forming a holistic understanding of human–AI interpretive co-creation.

References

- [1] Adobe Inc. 2025. *Adobe Premiere Pro*. <https://www.adobe.com/products/premiere.html> Accessed: 2025-11-05.
- [2] Torin Anderson and Shuo Niu. 2025. Making AI-Enhanced Videos: Analyzing Generative AI Use Cases in YouTube Content Creation. In *Proceedings of the Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*. 1–7.
- [3] Ahmet Yasin Aytar, Kemal Kilic, and Kamer Kaya. 2024. A retrieval-augmented generation framework for academic literature navigation in data science. *arXiv preprint arXiv:2412.15404* (2024).
- [4] Ava Bartolome and Shuo Niu. 2023. A literature review of video-sharing platform research in HCI. In *Proceedings of the 2023 CHI conference on human factors in computing systems*. 1–20.
- [5] Aadit Barua, Karim Benharrah, Meng Chen, Mina Huh, and Amy Pavel. 2025. Lotus: Creating short videos from long videos with abstractive and extractive summarization. In *Proceedings of the 30th International Conference on Intelligent User Interfaces*. 967–981.
- [6] Bytedance Ltd. 2025. *CapCut Video Editing Tools*. <https://www.capcut.com/tools> Accessed: 2025-11-05.
- [7] Yining Cao, Yiyi Huang, Anh Truong, Hijung Valentina Shin, and Haijun Xia. 2025. Compositional Structures as Substrates for Human-AI Co-creation Environment: A Design Approach and A Case Study. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–25.
- [8] Yan Chen, Andrés Monroy-Hernández, Ian Wehrman, Steve Oney, Walter S Lasecki, and Rajan Vaish. 2020. Sifter: A Hybrid Workflow for Theme-based Video Curation at Scale. In *Proceedings of the 2020 ACM International Conference on Interactive Media Experiences*. 65–73.
- [9] Zhi Chen. 2023. An Analysis of Narrative Language used in Film and Television Commentary Videos: Evidence from Bilibili.com. *International Journal of Education and Humanities* 1, 3 (2023), 18–23. Available on ResearchGate: https://www.researchgate.net/publication/366471140_An_Analysis_of_Narrative_Language_used_in_Film_and_Television_Commentary_Videos_Evidence_from_Bilibilicom.
- [10] Peggy Chi, Nathan Frey, Katrina Panovich, and Irfan Essa. 2021. Automatic instructional video creation from a markdown-formatted tutorial. In *The 34th Annual ACM Symposium on User Interface Software and Technology*. 677–690.
- [11] Zihan Ding, Junlong Chen, Per Ola Kristensson, Junxiao Shen, and Xinyi Wang. 2025. Prompt-Driven Agentic Video Editing System: Autonomous Comprehension of Long-Form, Story-Driven Media. *arXiv preprint arXiv:2509.16811* (2025).
- [12] Casey Fiesler and Amy S Bruckman. 2014. Remixers’ understandings of fair use online. In *Proceedings of the 17th ACM conference on Computer supported cooperative work & social computing*. 1023–1032.
- [13] Ohad Fried, Ayush Tewari, Michael Zollhöfer, Adam Finkelstein, Eli Shechtman, Dan B Goldman, Kyle Genova, Zeyu Jin, Christian Theobalt, and Maneesh Agrawala. 2019. Text-based editing of talking-head video. *ACM Transactions on Graphics (TOG)* 38, 4 (2019), 1–14.
- [14] Xiaoyan Fu. 2025. Analysis of the Popular Trends in Movie Commentary Videos. *Communications in Humanities Research* (2025). Available on ResearchGate: https://www.researchgate.net/publication/387716847_Analysis_of_the_Popular_Trends_in_Movie_Commentary_Videos.
- [15] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yixin Dai, Jiawei Sun, Haofen Wang, and Haofen Wang. 2023. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997* 2, 1

- (2023).
- [16] Hugo Flores García, Oriol Nieto, Justin Salamon, Bryan Pardo, and Prem Seetharaman. 2025. Sketch2sound: Controllable audio generation via time-varying signals and sonic imitations. In *ICASSP 2025-2025 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 1–5.
 - [17] Alan Hanjalic and Li-Qun Xu. 2005. Affective video content representation and modeling. *IEEE transactions on multimedia* 7, 1 (2005), 143–154.
 - [18] Tianhao He, Andrija Stanković, Evangelos Niforatos, and Gerd Kortuem. 2025. DesignMinds: Enhancing Video-Based Design Ideation with a Vision-Language Model and a Context-Injected Large Language Model. In *Proceedings of the 7th ACM Conference on Conversational User Interfaces*. 1–15.
 - [19] Bernd Huber, Hijung Valentina Shin, Bryan Russell, Oliver Wang, and Gautham J Mysore. 2019. B-script: Transcript-based b-roll video editing with recommendations. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*. 1–11.
 - [20] Emelia May Hughes, Renee Wang, Prerna Juneja, Tony W Li, Tanushree Mitra, and Amy X Zhang. 2024. Viblio: Introducing credibility signals and citations to video-sharing platforms. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. 1–20.
 - [21] Mina Huh, Ding Li, Kim Pimmel, Hijung Valentina Shin, Amy Pavel, and Mira Dontcheva. 2025. VideoDiff: Human-AI Video Co-Creation with Alternatives. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. 1–19.
 - [22] Jinkyu Jang, Jinwook Kim, Hyeonsik Shin, Hajung Aum, and Jinwoo Kim. 2016. Effects of temporal format of everyday video on narrative engagement and social interactivity. *Interacting with Computers* 28, 6 (2016), 718–736.
 - [23] Nicolas Jonason, Luca Casini, Carl Thomé, and Bob LT Sturm. 2023. Retrieval augmented generation of symbolic music with llms. *arXiv preprint arXiv:2311.10384* (2023).
 - [24] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in neural information processing systems* 33 (2020), 9459–9474.
 - [25] Yuzhi Li, Haojun Xu, and Feng Tian. 2025. Shot sequence ordering for video editing: Benchmarks, metrics, and cinematology-inspired computing methods. *arXiv preprint arXiv:2503.17975* (2025).
 - [26] Zhi Li, Sicheng Huai, and Yang Lin. 2024. Three minutes of fame: Effects of movie commentary video style on users' liking and commenting intentions. (2024).
 - [27] Jeba Rezwana and Mary Lou Maher. 2023. Designing creative AI partners with COFI: A framework for modeling interaction in human-AI co-creative systems. *ACM Transactions on Computer-Human Interaction* 30, 5 (2023), 1–28.
 - [28] Guy Schofield, Tom Bartindale, and Peter Wright. 2015. Bootlegger: turning fans into film crew. In *Proceedings of the 33rd annual ACM conference on human factors in computing systems*. 767–776.
 - [29] Renato Nureña Siles and Victor Omar Vite León. 2025. The Role of Film Critic YouTubers in the Development of Audiovisual Competencies. The Case of 3 Latin American YouTubers. In *International Conference on Communication and Applied Technologies 2025 (ICOMTA 2025)*. Atlantis Press, 466–478.
 - [30] Nikhil Singh, Guillermo Bernal, Daria Savchenko, and Elena L Glassman. 2023. Where to hide a stolen elephant: Leaps in creative writing with multimodal machine intelligence. *ACM Transactions on Computer-Human Interaction* 30, 5 (2023), 1–57.
 - [31] Erik Stolterman and Mikael Wiberg. 2010. Concept-driven interaction design research. *Human-Computer Interaction* 25, 2 (2010), 95–118.
 - [32] Blair Subbaraman, Shenna Shim, and Nadya Peek. 2023. Forking a sketch: How the openprocessing community uses remixing to collect, annotate, tune, and extend creative code. In *Proceedings of the 2023 ACM Designing Interactive Systems Conference*. 326–342.
 - [33] Jiatian Sun, Longxiulin Deng, Triantafyllos Afouras, Andrew Owens, and Abe Davis. 2023. Eventfulness for interactive video alignment. *ACM Transactions on Graphics (TOG)* 42, 4 (2023), 1–10.
 - [34] Yicheng Tao, Yao Qin, and Yepang Liu. 2025. Retrieval-Augmented Code Generation: A Survey with Focus on Repository-Level Approaches. *arXiv preprint arXiv:2510.04905* (2025).
 - [35] Yang Tian, Fan Liu, Jingyuan Zhang, Wei Bi, Yupeng Hu, and Liqiang Nie. 2025. Open Multimodal Retrieval-Augmented Factual Image Generation. *arXiv preprint arXiv:2510.22521* (2025).
 - [36] Bekzat Tilekbay, Saelyne Yang, Michal Adam Lewkowicz, Alex Suryapranata, and Juho Kim. 2024. Expressedit: Video editing with natural language and sketching. In *Proceedings of the 29th International Conference on Intelligent User Interfaces*. 515–536.
 - [37] Anh Truong, Floraine Berthouzoz, Wilmot Li, and Maneesh Agrawala. 2016. Quickcut: An interactive tool for editing narrated video. In *Proceedings of the 29th Annual Symposium on User Interface Software and Technology*. 497–507.
 - [38] Miao Wang, Guo-Wei Yang, Shi-Min Hu, Shing-Tung Yau, Ariel Shamir, et al. 2019. Write-a-video: computational video montage from themed text. *ACM Trans. Graph.* 38, 6 (2019), 177–1.

- [39] Sitong Wang, Samia Menon, Tao Long, Keren Henderson, Dingzeyu Li, Kevin Crowston, Mark Hansen, Jeffrey V Nickerson, and Lydia B Chilton. 2024. ReelFramer: Human-AI co-creation for news-to-video translation. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. 1–20.
- [40] Sitong Wang, Zheng Ning, Anh Truong, Mira Dontcheva, Dingzeyu Li, and Lydia B Chilton. 2024. PodReels: Human-AI Co-Creation of Video Podcast Teasers. In *Proceedings of the 2024 ACM Designing Interactive Systems Conference*. 958–974.
- [41] Zhihua Wen, Zhiliang Tian, Wei Wu, Yuxin Yang, Yanqi Shi, Zhen Huang, and Dongsheng Li. 2023. Grove: a retrieval-augmented complex story generation framework with a forest of evidence. *arXiv preprint arXiv:2310.05388* (2023).
- [42] Zhikun Wu, Thomas Weber, and Florian Müller. 2025. One does not simply meme alone: Evaluating co-creativity between LLMs and humans in the generation of humor. In *Proceedings of the 30th International Conference on Intelligent User Interfaces*. 1082–1092.
- [43] Yu Xiong, Fabian Caba Heilbron, and Dahua Lin. 2022. Transcript to video: Efficient clip sequencing from texts. In *Proceedings of the 30th ACM International Conference on Multimedia*. 5407–5416.
- [44] Guoxing Yang, Haoyu Lu, Zelong Sun, and Zhiwu Lu. 2023. Shot retrieval and assembly with text script for video montage generation. In *Proceedings of the 2023 ACM International Conference on Multimedia Retrieval*. 298–306.
- [45] Yun Yi, Yunkang Zhou, Tinghua Wang, and Jin Zhou. 2025. Advances in Video Emotion Recognition: Challenges and Trends. *Sensors* 25, 12 (2025), 3615.
- [46] Yiyi Yin and Anthony Fung. 2017. Youth online cultural participation and Bilibili: An alternative form of democracy in China? In *Digital media integration for participatory democracy*. IGI Global Scientific Publishing, 130–154.
- [47] Ann Yuan, Andy Coenen, Emily Reif, and Daphne Ippolito. 2022. Wordcraft: story writing with large language models. In *Proceedings of the 27th International Conference on Intelligent User Interfaces*. 841–852.
- [48] Shuning Zhang, Hui Wang, and Xin Yi. 2025. Exploring Collaboration Patterns and Strategies in Human-AI Co-creation through the Lens of Agency: A Scoping Review of the Top-tier HCI Literature. *Proceedings of the ACM on Human-Computer Interaction* 9, 7 (2025), 1–43.
- [49] Chongjun Zhong, Jiaxing Yu, Yingping Cao, Songruoyao Wu, Wenqi Wu, and Kejun Zhang. 2024. SoundScape: A Human-AI Co-Creation System Making Your Memories Heard. *arXiv preprint arXiv:2410.08136* (2024).
- [50] John Zimmerman, Jodi Forlizzi, and Shelley Evenson. 2007. Research through design as a method for interaction design research in HCI. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. ACM, 493–502.
- [51] Qiaoling Zou, Wanyu Zheng, Zishun Su, Li Zhang, Ziqing Zhuo, and Dongning Li. 2025. The neural impact of editing on viewer narrative cognition in virtual reality films: eye-tracking insights into neural mechanisms. *Frontiers in Psychology* 16 (2025), 1584250.